

# INTRO TO DEEP LEARNING FOCUSING ON TOOLS



**Facebook.com/knowlexon**

**<https://www.linkedin.com/company/knowlexon>**

**info@knowlexon.com**

**Biswa G Singh**  
**Acknowledgement: Subrat Panda**

# SPEAKER SHORT BIO


- MS (2010) – Clemson University
- Synopsys (EDA), IBM (CPU), ARM, Qualcomm, HP, AMD
- Lead ML Engineer at Capillary Technologies (We are Hiring)
- Deep Learning Performance on CPU/GPU
- Co-Founded IDLI (for social good) with Prof. Amit Sethi (IIT Bombay), **Subrat Panda (Capillary AI head, IIT KGP) and Jacob Minz (Synopsys)**
- <https://www.facebook.com/groups/idliai/> (Consider Joining)
- Linked In - <https://www.linkedin.com/in/biswagsingh/>
- **Mentor** to knowlexon Innovation and technology Pvt Ltd(knowlexon.com)



“AI IS THE NEW  
ELECTRICITY.”

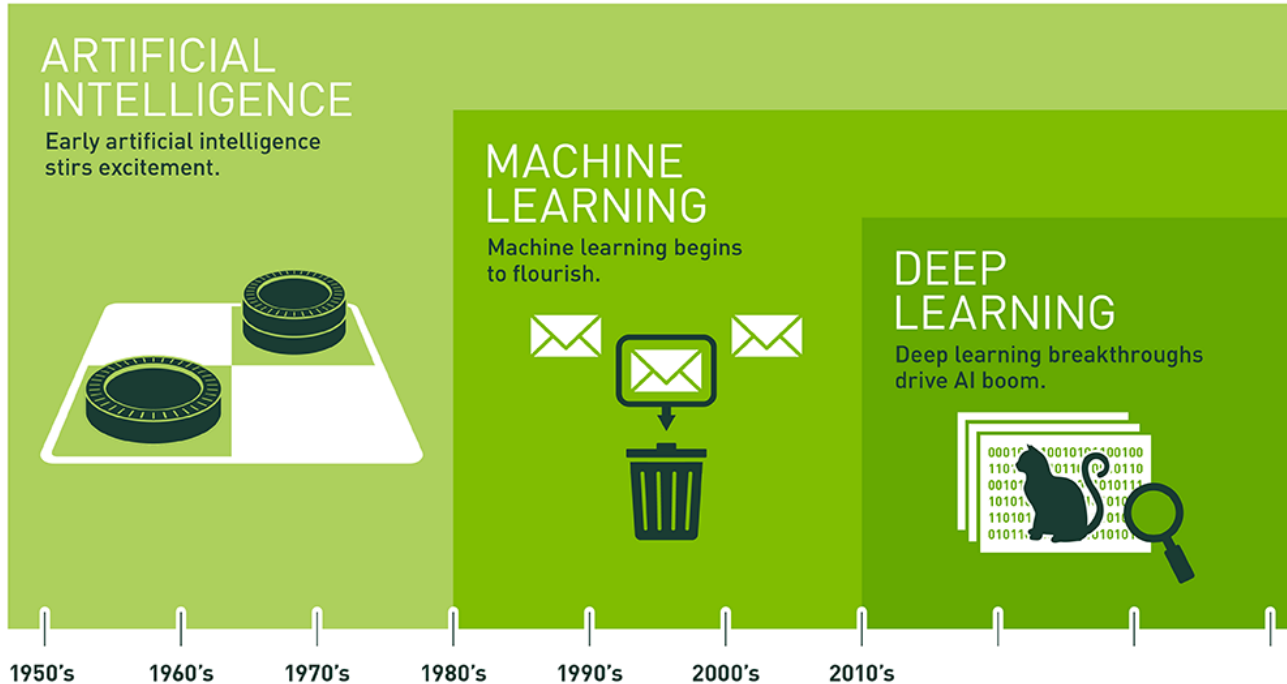
“Just as electricity transformed almost everything 100 years ago, today I actually have a hard time thinking of an industry that I don’t think AI will transform in the next several years.”

**Andrew Ng**



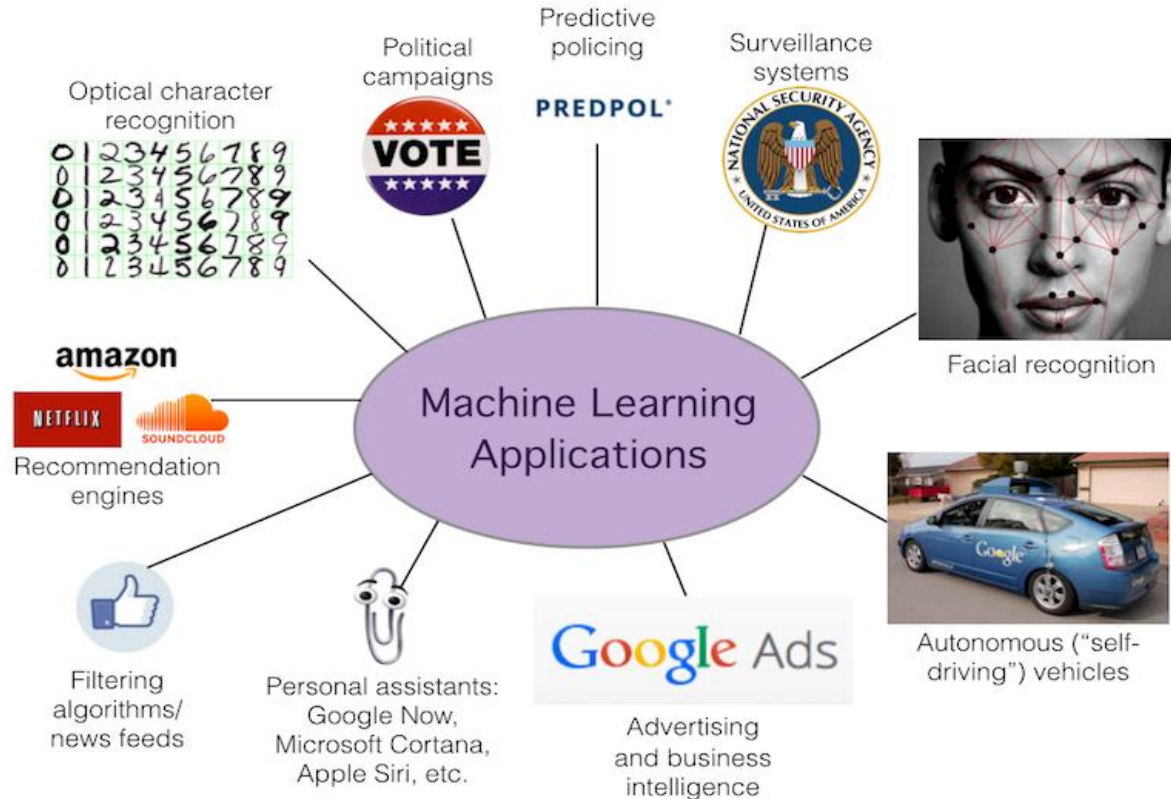
# Preface

- Artificial intelligence is already part of our everyday lives.





Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

# Application of AI, Machine Learning and Deep Learning



# Artificial Neural Networks

- The idea of ANNs is based on the belief that working of human brain by making the right connections, can be imitated using silicon and wires as living **neurons and dendrites**.
- **A Single Neuron:** The basic unit of computation in a neural network is the **neuron**, often called a **node** or **unit**.



- The **function  $f$**  is non-linear and is called the **Activation Function**.

# Activation Function

- **Sigmoid:** takes a real-valued input and squashes it to range between 0 and 1.


$$\sigma(x) = 1 / (1 + \exp(-x))$$

- **tanh:** takes a real-valued input and squashes it to the range  $[-1, 1]$


$$\tanh(x) = 2\sigma(2x) - 1$$

- **ReLU:** ReLU stands for Rectified Linear Unit. It takes a real-valued input and thresholds it at zero (replaces negative values with zero)


$$f(x) = \max(0, x)$$



Sigmoid




tanh




ReLU


# Neural Network Intuition (single layer)




$X_1$	NOT $X_1$	$(1-2*X_1)$	$a$
0	1	1	1
1	0	-1	0



$X_1$	$X_2$	$X_1$ OR $X_2$	$(-0.5+X_1+X_2)$	$a$
0	0	0	-0.5	0
0	1	1	0.5	1
1	0	1	0.5	1
1	1	1	1.5	1



$X_1$	$X_2$	$X_1$ AND $X_2$	$(-1.5+X_1+X_2)$	$a$
0	0	0	-1.5	0
0	1	0	-0.5	0
1	0	0	-0.5	0
1	1	1	0.5	1



$X_1$	$X_2$	$X_1'$ AND $X_2'$	$(0.5-X_1-X_2)$	$a$
0	0	1	0.5	1
0	1	0	-0.5	0
1	0	0	-0.5	0
1	1	0	-1.5	0




# Neural Network Intuition (Multiple Layer layer)




info@knowlexon.com

- Multi Layer Neural network is capable of learning complex functions.
- Lets consider XNOR operation.

- CASE1:  $X_1 \text{ XNOR } X_2 = (A'.B') + (A.B)$




- CASE2:  $X_1 \text{ XNOR } X_2 = \text{NOT} [(A+B).(A'+B')]$




NN representation

NN representation = ?


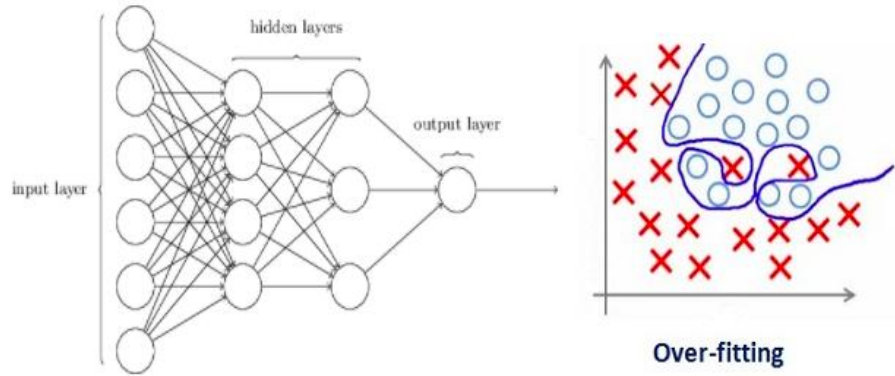


# Back-Propagation

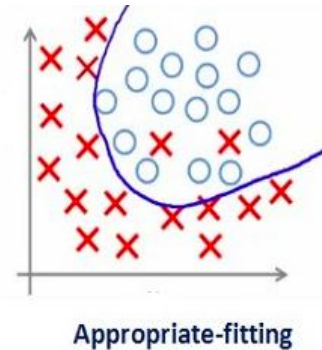
- Back-propagation (BP) algorithm works by determining the loss (or error) at the output and then propagating it back into the network.
- The weights are updated to minimize the error resulting from each neuron.



# Regularization: Dropout



- At every iteration, it randomly selects some nodes and removes them along with all of their incoming and outgoing connections
- We need to choose the dropout parameter such that we get the appropriate fitting



# Building a Deep Learning Application

- Make sure that you actually need Deep Learning - don't just jump into it.
- You need to have DATA (Labelled) - curated.
- GIGO is the rule on using data.
- Choose the right algorithm/network/framework for your Use case.
- Find out if you can use **Transfer Learning** - if yes go ahead try it out.
- If you don't have data - Prepare a data capture strategy (buy it. Collect it etc) - and use API based service of known providers if you can use it.
- Do not reinvent the wheel of trying to build a network from scratch.
- Play with hyper-parameters of known models/framework etc.
- Some cases an ensemble is good enough for the problem.
- Look at Kaggle, Arxiv and other data science competition platforms

# Different DL Frameworks



theano



# Underneath



# Basic Components: DL Frameworks

- Tensors
- Operations
- Computational Graph: Combine multiple operations
- Auto-differentiation Module
- GPU support with FP operation (cuBLAS etc)

# What is a Tensor?

- Tensors are multidimensional arrays
- This is the standard way of representing data in Tensorflow

't'
'e'
'n'
's'
'o'
'r'

*Tensor of  
dimension[1]*

3	1	4	1
5	9	2	6
5	3	5	8
9	7	9	3
2	3	8	4
6	2	6	4

*Tensor of  
dimensions[2]*

2	1	2	8	1	8
2	8	5	0	4	5
2	5	3	0	2	8
7	4	7	1	5	6

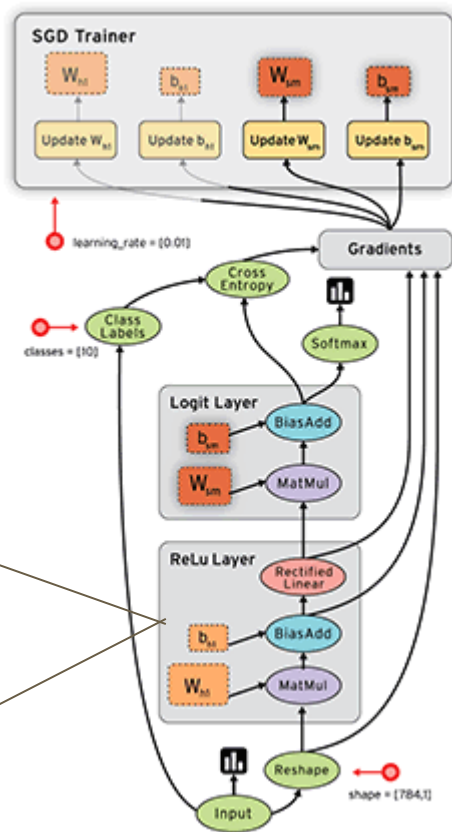

*Tensor of  
dimensions[3]*



# TensorFlow, a DL Library

- 1) A deep learning library open sourced by google
- 2) Created for task with heavy numeric computation
- 3) Based on dataflow graph

$$h = \text{ReLU}(Wx + b)$$



# Tensorflow code


```
import tensorflow as tf

b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))

x = tf.placeholder(tf.float32, (100, 784))

h = tf.nn.relu(tf.matmul(x, W) + b)
```

$$h = \text{ReLU}(Wx + b)$$



# Tensorflow code: Getting the output

`sess.run(fetches, feeds)`

**Fetches:** List of graph nodes.

To fetch the the output of the operation execute the graph in `run()`.

**Feeds:**  
Feed data to the graph

```
import numpy as np
import tensorflow as tf

b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100),
                                  -1, 1))

x = tf.placeholder(tf.float32, (100, 784))
h = tf.nn.relu(tf.matmul(x, W) + b)

sess = tf.Session()
sess.run(tf.initialize_all_variables())
sess.run(h, {x: np.random.random(100, 784)})
```

# So What we covered so far?

We first built a **graph** using **variables**, **placeholders** and **operations** on them


We then deployed the graph **onto a session**, which is the **execution environment**

Next we will see **how to train the model**

# Handwritten Digit classification MNIST



28x28  
pixels



Info@knowlexon.com


0	4	1	2	1	3	1	4	3
5	3	6	1	7	2	8	6	9
0	9	1	1	2	4	3	2	7
8	6	9	0	5	6	0	7	6
8	1	9	3	9	8	5	9	3
0	7	4	9	8	0	9	4	1
4	6	0	4	5	6	1	0	1
7	1	6	3	0	2	1	1	7
0	2	6	7	8	3	9	0	4
7	4	6	8	0	7	8	3	1

weighted sum of all  
pixels + bias

$$\text{softmax}(L_n) = \frac{e^{L_n}}{\|e^L\|}$$

neuron outputs

# Matrix Notation



# Error Metric

0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	1	0	0	0

actual probabilities, "one-hot" encoded

Cross entropy:  $-\sum Y_i' \cdot \log(Y_i)$

computed probabilities

0.1	0.2	0.1	0.3	0.2	0.1	0.9	0.2	0.1	0.1
0	1	2	3	4	5	6	7	8	9

this is a "6"

# Training

```
optimizer = tf.train.GradientDescentOptimizer(0.003)  
train_step = optimizer.minimize(cross_entropy)
```




# Dropout


A regularization technique works great in Deep Neural network

Randomly Drops certain % of nodes in training to reduce overfitting

An efficient way of doing model averaging as we randomly drop nodes for each training example




TRAINING  
 $p_{\text{Keep}}=0.75$



EVALUATION  
 $p_{\text{Keep}}=1$

# DEMO



# Convolutional Neural Network (CNN)

info@knowlexon.com


Problem With Fully connected network on images:

Losing shape information when we are flattening the image into a single array


Visual cortex system of brain doesn't work like fully connected network.

**Huble and Wiesel** experiment on cat reveals

- 1) **Local connection**
- 2) **Hierarchical layer**
- 3) **Spatial Invariance : Any size, rotation, shift**





# CNN to rescue




[LeNet-5, LeCun 1998]

# Convolution layer



$$\begin{array}{l}
 W[4, 4, 3] \\
 W_2[4, 4, 3]
 \end{array}
 \left|
 \begin{array}{l}
 W[4, 4, 3, 2] \\
 \text{filter size} \quad \text{input channels} \quad \text{output channels}
 \end{array}
 \right.$$

# Conv Net on MNIST




# Pooling

The idea of pooling in convolutional neural networks is to do two things:

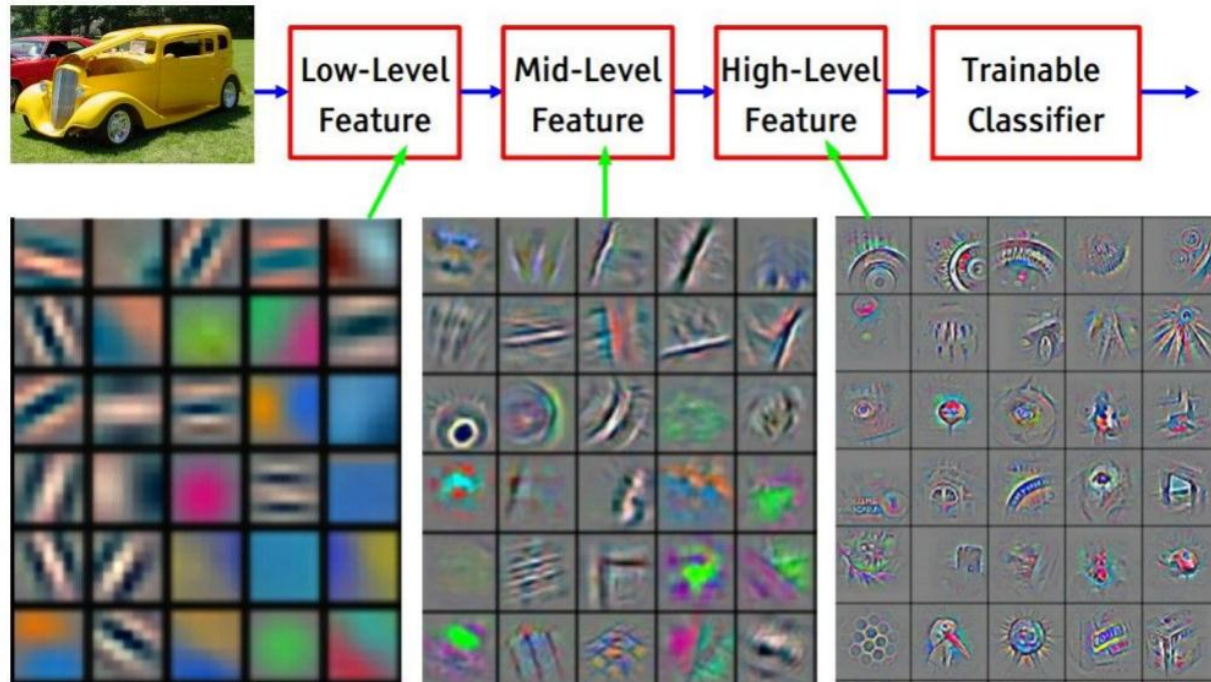
- Reduce the number of parameters in your network (pooling is also called “down-sampling” for this reason)
- To make feature detection more robust by making it more impervious to scale and orientation changes

Max Pooling

Average Pooling



# Feature extraction



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]



# Pytorch vs Tensorflow

- Tensorflow is widely adopted but pytorch picking up
- Dynamic vs static graph
- Tensorboard is better than pytorch visualization
- Plain tensorflow looks pretty much like a library
- Abstraction is better in pytorch, even data parallelism
- Tf.contrib, keras to rescue

# PYTORCH Framework Pieces

1. torch: a general purpose array library similar to Numpy that can do computations on GPU when the **tensor type is cast to (torch.cuda.FloatTensor)**
2. torch.autograd: a package for building a computational graph and automatically obtaining gradients
3. torch.nn: a neural net library with common layers and cost functions
4. torch.optim: an optimization package with common optimization algorithms like SGD, Adam, etc

**DEMO**

# Acknowledge

- Materials have been taken from
- Martin Gorner Tensorflow Tutorial
- Stanford CS224d Tensorflow Tutorial
- Stanford CS231n course material
- Analytics Vidya article
- Many more PPTs available on the internet from Universities and Good Samaritans
- [www.tensorflow.org](http://www.tensorflow.org)

# For more details -

[www.Facebook.com/knowlexon](https://www.facebook.com/knowlexon)

<https://www.linkedin.com/company/knowlexon>

[www.Knowlexon.com](http://www.Knowlexon.com)

[info@knowlexon.com](mailto:info@knowlexon.com)